

Practical privacy using homomorphic encryption – a myth or reality?

Anirban Basu

Department of Electrical Engineering
Faculty of Engineering
Tokai University (Japan)

SecureCloud 2012
May 9, 2012
Frankfurt am Main, Germany

The magic of homomorphic encryption

1 Homomorphic encryption

- What and why?
- Application in recommender systems

2 Privacy preserving collaborative filtering (PPCF)

- Collaborative filtering (CF), briefly
- Privacy preserving Slope One


3 On the SaaS cloud

- Google App Engine feasibility
- An alternative approach


4 Tail piece

- Conclusions and future avenues
- Question time!

Compute blindly?



Compute blindly?



Compute blindly?

$$(m_1 \oplus m_2)' = m_1' \otimes m_2'$$

m_1 and m_2 : plaintext messages; x' : ciphertext of x ; \oplus and \otimes : two operations that satisfy this **homomorphic** relation.

Homomorphic cryptosystems

Depending on \oplus and \otimes , we have:

- additive (e.g., Paillier, Damgård-Jurik, Elliptic Curve ElGamal),
- multiplicative (e.g., RSA, ElGamal), and
- fully homomorphic cryptosystems (Craig Gentry et al.).
- too good to be true?

Homomorphic cryptosystems

Depending on \oplus and \otimes , we have:

- additive (e.g., Paillier, Damgård-Jurik, Elliptic Curve ElGamal),
- **multiplicative (e.g., RSA, ElGamal), and**
- fully homomorphic cryptosystems (Craig Gentry et al.).
- too good to be true?
- or, practically feasible?

Homomorphic cryptosystems

Depending on \oplus and \otimes , we have:

- additive (e.g., Paillier, Damgård-Jurik, Elliptic Curve ElGamal),
 - multiplicative (e.g., RSA, ElGamal), and
 - **fully homomorphic cryptosystems (Craig Gentry et al.).**
- too good to be true?
- or, practically feasible?

Homomorphic cryptosystems

Depending on \oplus and \otimes , we have:

- additive (e.g., Paillier, Damgård-Jurik, Elliptic Curve ElGamal),
- multiplicative (e.g., RSA, ElGamal), and
- fully homomorphic cryptosystems (Craig Gentry et al.).

but, is this *magic*

- too good to be true?
- or, practically feasible?

Homomorphic cryptosystems

Depending on \oplus and \otimes , we have:

- additive (e.g., Paillier, Damgård-Jurik, Elliptic Curve ElGamal),
- multiplicative (e.g., RSA, ElGamal), and
- fully homomorphic cryptosystems (Craig Gentry et al.).

but, is this *magic*

- too good to be true?
- or, practically feasible?

Recommendation through collaborative filtering

- *‘People who have bought this have also bought these’ – recommendation.*
- Collaborative filtering (CF) – recommendation from opinions of the community.
- What about privacy in rating based collaborative filtering?

Recommendation through collaborative filtering

- *‘People who have bought this have also bought these’* – recommendation.
- Collaborative filtering (CF) – recommendation from opinions of the community.
- What about privacy in rating based collaborative filtering?

Recommendation through collaborative filtering

- *'People who have bought this have also bought these'* – recommendation.
- Collaborative filtering (CF) – recommendation from opinions of the community.
- What about privacy in rating based collaborative filtering?

Privacy and recommender systems

- Can *someone* (the cloud?) compute CF for users?
- ... and do so blindly?
- Privacy preserving collaborative filtering (PPCF) for the Software-as-a-Service cloud.

Privacy and recommender systems

- Can *someone* (the cloud?) compute CF for users?
- ...and do so blindly?
- Privacy preserving collaborative filtering (PPCF) for the Software-as-a-Service cloud.

Privacy and recommender systems

- Can *someone* (the cloud?) compute CF for users?
- ... and do so blindly?
- Privacy preserving collaborative filtering (PPCF) for the Software-as-a-Service cloud.

Application scenario

- 1 Homomorphic encryption
 - What and why?
 - Application in recommender systems
- 2 Privacy preserving collaborative filtering (PPCF)
 - Collaborative filtering (CF), briefly
 - Privacy preserving Slope One
- 3 On the SaaS cloud
 - Google App Engine feasibility
 - An alternative approach
- 4 Tail piece
 - Conclusions and future avenues
 - Question time!

CF – a form of recommendation

User-item rating data like this¹:

	Canon 7D	Leica M9	Nikon D7000	...	Olympus OM-D
Alice	5	4	-	...	3
Bob	3	5	2	...	-
Carol	-	?	4	...	3
Dave	4	3	-	...	-

- The objective is to find a rating for Leica M9 for Carol.
- CF – a well-known recommendation technique, based on the preferences of the community.

¹“-” indicates the absence of a rating.

CF – a form of recommendation

User-item rating data like this¹:

	Canon 7D	Leica M9	Nikon D7000	...	Olympus OM-D
Alice	5	4	-	...	3
Bob	3	5	2	...	-
Carol	-	?	4	...	3
Dave	4	3	-	...	-

- The objective is to find a rating for Leica M9 for Carol.
- CF – a well-known recommendation technique, based on the preferences of the community.

¹“-” indicates the absence of a rating.

Slope One predictors

Based on:

Lemire, D., Maclachlan, A. 2005. *Slope one predictors for online rating-based collaborative filtering*. In: Society for Industrial Mathematics.

Slope One predictors

- CF predictors of the form $f(x) = x + b$, hence “slope one”.
- Simple and efficient (compared with SVD, PMCC, Cosine).

Slope One predictors

- CF predictors of the form $f(x) = x + b$, hence “slope one”.
- Simple and efficient (compared with SVD, PMCC, Cosine).

Slope One predictors




Figure: The general CF problem.

CF and Slope One

Pre-compute:

- **Deviation matrix** Δ : deviation of ratings of an item pair by the same user; dimension: $n \times n$.
- **Cardinality matrix** ϕ : number of co-existing ratings by the same user of an item pair; dimension same as Δ .

CF and Slope One

Pre-compute:

- **Deviation matrix** Δ : deviation of ratings of an item pair by the same user; dimension: $n \times n$.
- **Cardinality matrix** ϕ : number of co-existing ratings by the same user of an item pair; dimension same as Δ .

CF and Slope One




Figure: Slope One pre-computation creates a ‘model’ which is used for prediction.

The weighted Slope One predictor

■ Average deviation:

$$\overline{\delta_{a,b}} = \frac{\Delta_{a,b}}{\phi_{a,b}} = \frac{\sum_i \delta_{i,a,b}}{\phi_{a,b}} = \frac{\sum_i (r_{i,a} - r_{i,b})}{\phi_{a,b}}$$

$\phi_{a,b}$: the number of the users who have rated both items;
 $\delta_{i,a,b} = r_{i,a} - r_{i,b}$: the deviation of the ratings of item a from that of item b both given by user i .

■ The *weighted* Slope One predictor:

$$r_{u,x} = \frac{\sum_{a|a \neq x} (\overline{\delta_{x,a}} + r_{u,a}) \phi_{x,a}}{\sum_{a|a \neq x} \phi_{x,a}} = \frac{\sum_{a|a \neq x} (\Delta_{x,a} + r_{u,a} \phi_{x,a})}{\sum_{a|a \neq x} \phi_{x,a}}$$

The weighted Slope One predictor

■ Average deviation:

$$\overline{\delta_{a,b}} = \frac{\Delta_{a,b}}{\phi_{a,b}} = \frac{\sum_i \delta_{i,a,b}}{\phi_{a,b}} = \frac{\sum_i (r_{i,a} - r_{i,b})}{\phi_{a,b}}$$

$\phi_{a,b}$: the number of the users who have rated both items;
 $\delta_{i,a,b} = r_{i,a} - r_{i,b}$: the deviation of the ratings of item a from that of item b both given by user i .

■ The *weighted* Slope One predictor:

$$r_{u,x} = \frac{\sum_{a|a \neq x} (\overline{\delta_{x,a}} + r_{u,a}) \phi_{x,a}}{\sum_{a|a \neq x} \phi_{x,a}} = \frac{\sum_{a|a \neq x} (\Delta_{x,a} + r_{u,a} \phi_{x,a})}{\sum_{a|a \neq x} \phi_{x,a}}$$

Preserving privacy with Slope One CF

- An *additively homomorphic cryptosystem* – the Damgård-Jurik cryptosystem, defining

- homomorphic addition:

$$\mathcal{E}(m_1 + m_2) = \mathcal{E}(m_1) \cdot \mathcal{E}(m_2)$$

- and homomorphic multiplication:

$$\mathcal{E}(m_1 \cdot \pi) = \mathcal{E}(m_1)^\pi$$

m_1 and m_2 are plaintexts and π is a plaintext multiplicand.

Preserving privacy with Slope One CF

- An *additively homomorphic cryptosystem* – the Damgård-Jurik cryptosystem, defining

- **homomorphic addition:**

$$\mathcal{E}(m_1 + m_2) = \mathcal{E}(m_1) \cdot \mathcal{E}(m_2)$$

- and homomorphic multiplication:

$$\mathcal{E}(m_1 \cdot \pi) = \mathcal{E}(m_1)^\pi$$

m_1 and m_2 are plaintexts and π is a plaintext multiplicand.

Preserving privacy with Slope One CF

- An *additively homomorphic cryptosystem* – the Damgård-Jurik cryptosystem, defining
 - homomorphic addition:

$$\mathcal{E}(m_1 + m_2) = \mathcal{E}(m_1) \cdot \mathcal{E}(m_2)$$

- and homomorphic multiplication:

$$\mathcal{E}(m_1 \cdot \pi) = \mathcal{E}(m_1)^\pi$$

m_1 and m_2 are plaintexts and π is a plaintext multiplicand.

Privacy preserving weighted Slope One predictors

- The numerator of the prediction equation:

$$\sum_{a|a \neq x} (\Delta_{x,a} + r_{u,a} \phi_{x,a}) = \mathcal{D} \left(\prod_{a|a \neq x} (\mathcal{E}(\Delta_{x,a}) (\mathcal{E}(r_{u,a})^{\phi_{x,a}})) \right)$$

- The final prediction:

$$r_{u,x} = \frac{\mathcal{D}(\prod_{a|a \neq x} (\mathcal{E}(\Delta_{x,a}) (\mathcal{E}(r_{u,a})^{\phi_{x,a}})))}{\sum_{a|a \neq x} \phi_{x,a}}$$

Privacy preserving weighted Slope One predictors

- The numerator of the prediction equation:

$$\sum_{a|a \neq x} (\Delta_{x,a} + r_{u,a} \phi_{x,a}) = \mathcal{D} \left(\prod_{a|a \neq x} (\mathcal{E}(\Delta_{x,a}) (\mathcal{E}(r_{u,a})^{\phi_{x,a}})) \right)$$

- The final prediction:

$$r_{u,x} = \frac{\mathcal{D}(\prod_{a|a \neq x} (\mathcal{E}(\Delta_{x,a}) (\mathcal{E}(r_{u,a})^{\phi_{x,a}})))}{\sum_{a|a \neq x} \phi_{x,a}}$$

Privacy preserving weighted Slope One predictors

- Pre-computed: $\mathcal{E}(\Delta)$ and ϕ .
- Encrypted query result decrypted using threshold decryption keys.
- Query answered by collaborating sites.
- Supports horizontal and vertical partitioning, using secure scalar product (Vaidya and Clifton).

Privacy preserving weighted Slope One predictors

- Pre-computed: $\mathcal{E}(\Delta)$ and ϕ .
- Encrypted query result decrypted using threshold decryption keys.
- Query answered by collaborating sites.
- Supports horizontal and vertical partitioning, using secure scalar product (Vaidya and Clifton).

Privacy preserving weighted Slope One predictors

- Pre-computed: $\mathcal{E}(\Delta)$ and ϕ .
- Encrypted query result decrypted using threshold decryption keys.
- Query answered by collaborating sites.
- Supports horizontal and vertical partitioning, using secure scalar product (Vaidya and Clifton).

Privacy preserving weighted Slope One predictors

- Pre-computed: $\mathcal{E}(\Delta)$ and ϕ .
- Encrypted query result decrypted using threshold decryption keys.
- Query answered by collaborating sites.
- Supports horizontal and vertical partitioning, using secure scalar product (Vaidya and Clifton).

Some results

- **Hardware:** 2.53GHz Intel Core 2 Duo (dual core) processor, 8GB RAM, Mac OS X 10.6.7 and 64-bit Java 1.6.
- **Implementation:** Single server, single partition, MovieLens 100K dataset.

Some results

- **Hardware:** 2.53GHz Intel Core 2 Duo (dual core) processor, 8GB RAM, Mac OS X 10.6.7 and 64-bit Java 1.6.
- **Implementation:** Single server, single partition, MovieLens 100K dataset.

Key size	Mean PC ^a	Total PC	Prediction
512 bits	30s	2h	500ms
1024 bits	90s	6h	1.5s
2048 bits	270s	18h	4.5s

^aPC: pre-computation

Further reading

JISIS 2011 paper:

Basu, A., Vaidya, J., Kikuchi, H. 2011. *Efficient privacy-preserving collaborative filtering based on the weighted Slope One predictor*. In: Journal of Internet Services and Information Security (special edition).

Try this on the cloud

- 1 Homomorphic encryption
 - What and why?
 - Application in recommender systems
- 2 Privacy preserving collaborative filtering (PPCF)
 - Collaborative filtering (CF), briefly
 - Privacy preserving Slope One
- 3 On the SaaS cloud
 - Google App Engine feasibility
 - An alternative approach
- 4 Tail piece
 - Conclusions and future avenues
 - Question time!

Google App Engine for Java (GAE/J)

- **A Software-as-a-Service (SaaS) engine.**
 - On-demand transparent scalability with low costs, including a daily free quota.
 - Java servlet based computation model; allows for batch computations using task queues and cron.

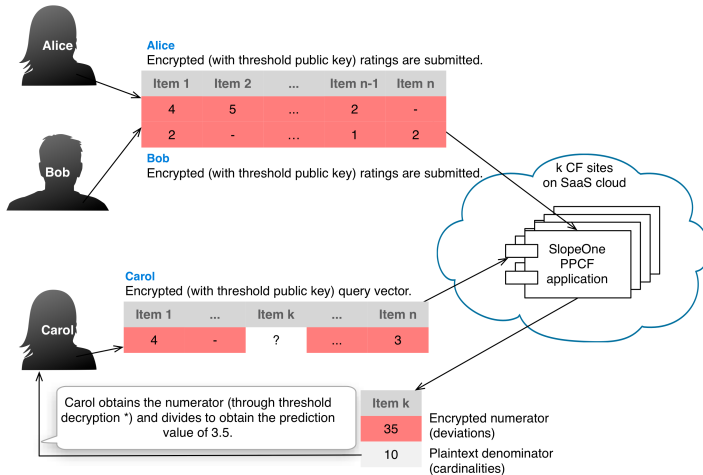
Google App Engine for Java (GAE/J)

- A Software-as-a-Service (SaaS) engine.
- On-demand transparent scalability with low costs, including a daily free quota.
- Java servlet based computation model; allows for batch computations using task queues and cron.

Google App Engine for Java (GAE/J)

- A Software-as-a-Service (SaaS) engine.
- On-demand transparent scalability with low costs, including a daily free quota.
- **Java servlet based computation model; allows for batch computations using task queues and cron.**

PPCF on the GAE/J



* Multiple PPCF sites compute the threshold decryption.

GAE/J challenges

- Servlet execution time limit (30s).
- High replication, slow access datastore.
- Lack of concurrency support (no threads!).
- Lack of control over resource allocation (bit better now with front-end instance classes).

GAE/J challenges

- Servlet execution time limit (30s).
- High replication, slow access datastore.
- Lack of concurrency support (no threads!).
- Lack of control over resource allocation (bit better now with front-end instance classes).

GAE/J challenges

- Servlet execution time limit (30s).
- High replication, slow access datastore.
- Lack of concurrency support (no threads!).
- Lack of control over resource allocation (bit better now with front-end instance classes).

GAE/J challenges

- Servlet execution time limit (30s).
- High replication, slow access datastore.
- Lack of concurrency support (no threads!).
- Lack of control over resource allocation (bit better now with front-end instance classes).

Further reading

ACM SAC 2012 paper:

Basu, A., Vaidya, J., Dimitrakos, T., Kikuchi, H. 2012. *Feasibility of a privacy preserving collaborative filtering scheme on the Google App Engine - a performance case study*. In Proc: 27th ACM Symposium on Applied Computing (SAC) Cloud Computing track, Trento.

An alternative PPCF on the GAE/J

- Identity anonymizer, plaintext deviation aggregation.
- No threshold decryption – user encrypts and decrypts.
- Prediction:

$$r_{u,x} = \frac{\mathcal{D}(\mathcal{E}(\sum_{a|a \neq x} \Delta_{x,a}) \prod_{a|a \neq x} (\mathcal{E}(r_{u,a})^{\phi_{x,a}}))}{\sum_{a|a \neq x} \phi_{x,a}}$$

An alternative PPCF on the GAE/J

- Identity anonymizer, plaintext deviation aggregation.
- No threshold decryption – user encrypts and decrypts.
- Prediction:

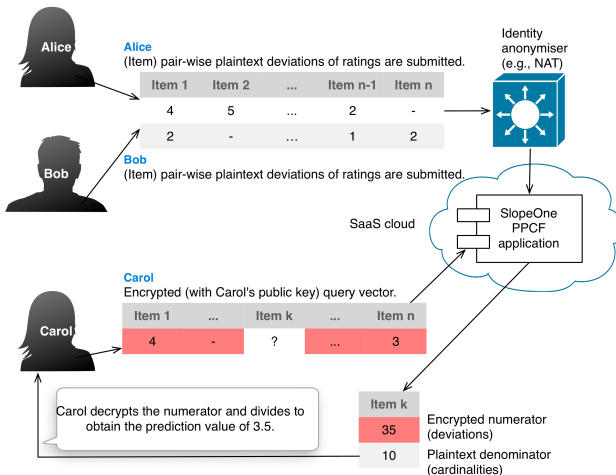
$$r_{u,x} = \frac{\mathcal{D}(\mathcal{E}(\sum_{a|a \neq x} \Delta_{x,a}) \prod_{a|a \neq x} (\mathcal{E}(r_{u,a})^{\phi_{x,a}}))}{\sum_{a|a \neq x} \phi_{x,a}}$$

An alternative PPCF on the GAE/J

- Identity anonymizer, plaintext deviation aggregation.
- No threshold decryption – user encrypts and decrypts.
- Prediction:

$$r_{u,x} = \frac{\mathcal{D}(\mathcal{E}(\sum_{a|a \neq x} \Delta_{x,a}) \prod_{a|a \neq x} (\mathcal{E}(r_{u,a})^{\phi_{x,a}}))}{\sum_{a|a \neq x} \phi_{x,a}}$$

An alternative PPCF on the GAE/J




Results on the Google App Engine (GAE/J)

Bit size ^a	Vector size ^b	Prediction time
1024	5	500ms
1024	10	650ms
2048	5	3800ms
2048	10	5000ms

^aPaillier cryptosystem modulus bit size, i.e. $|n|$.

^bSize of the encrypted rating query vector.

Results: GAE/J and Amazon Elastic Beanstalk



Further reading

IEEE Cloudcom 2011 paper:

Basu, A., Vaidya, J., Kikuchi, H., Dimitrakos, T. 2011.
Privacy-preserving collaborative filtering for the cloud. In Proc:
3rd IEEE International Conference on Cloud Computing
Technology and Science (Cloudcom), Athens.

Let's wrap up

- 1 Homomorphic encryption
 - What and why?
 - Application in recommender systems
- 2 Privacy preserving collaborative filtering (PPCF)
 - Collaborative filtering (CF), briefly
 - Privacy preserving Slope One
- 3 On the SaaS cloud
 - Google App Engine feasibility
 - An alternative approach
- 4 Tail piece
 - Conclusions and future avenues
 - Question time!

Conclusions

Privacy with homomorphic encryption – worth it?

It works well in certain scenarios albeit some, mainly, performance issues.

Conclusions

- Homomorphic encryption challenges: threshold decryption, computational complexity, supported operations.
- Application areas: privacy preserving data mining, graph problems e.g., resource allocation.
- Cloud applicability: performance, feasibility depends on type of application, platform.

Conclusions

- Homomorphic encryption challenges: threshold decryption, computational complexity, supported operations.
- Application areas: privacy preserving data mining, graph problems e.g., resource allocation.
- Cloud applicability: performance, feasibility depends on type of application, platform.

Conclusions

- Homomorphic encryption challenges: threshold decryption, computational complexity, supported operations.
- Application areas: privacy preserving data mining, graph problems e.g., resource allocation.
- Cloud applicability: performance, feasibility depends on type of application, platform.

Privacy, cloud, responsibility and the future

- Better performance – dedicated cryptographic hardware?
- Optimisation for computational complexity, power use – cloud infrastructure consumes power.
- Alternatives – open question?

Privacy, cloud, responsibility and the future

- Better performance – dedicated cryptographic hardware?
- Optimisation for computational complexity, power use – cloud infrastructure consumes power.
- Alternatives – open question?

Privacy, cloud, responsibility and the future

- Better performance – dedicated cryptographic hardware?
- Optimisation for computational complexity, power use – cloud infrastructure consumes power.
- Alternatives – open question?

Thank you for listening!

Any questions?